

CS 320: Principles of Programming Languages

Katie Casamento, Portland State University
Based on slides by Mark P. Jones, Portland State University, Winter 2017

Winter 2018
Week 3: Describing Syntax

Language = **syntax** + semantics

concrete syntax: the representation of a program text in its source form as a sequence of bits/bytes/characters/lines

abstract syntax: the representation of a program structure, independent of written form

syntax analysis: transformation from concrete syntax to abstract syntax

Syntax analysis

Usually a two-step process:

- *Tokenization* or *lexing* turns a character stream into a token stream
- *Parsing* turns a token stream into an abstract syntax structure

Syntax analysis


Character stream

```
"if y == 0 then Nothing else Just (x / y)"
```

→ Token stream

```
[IF, ID(y), ==, NUM(0), THEN, ID(Nothing), ELSE, ID(Just),  
L_PAREN, ID(x), /, ID(y), R_PAREN]
```

→ Abstract syntax tree



Formal languages

How do we describe syntax?

We want methods that are:

- clear, precise, and unambiguous
- expressive (e.g. finite descriptions of infinite languages)
- suitable for use in the implementation of syntax analysis tools (lexers, parsers, ...)

Formal languages provides such a foundation:

- *Regular languages* describe lexical syntax (grouping characters into tokens)
- *Context-free languages* describe more complex syntactic structure (parsing token streams into expressions)

Formal languages

- Pick a set of *symbols*, A , to be the *alphabet*
 - For lexical analysis, "symbols" are typically characters
 - For parsing, "symbols" are typically tokens
- The set of all finite strings of symbols in A is written A^*
- A *language* over A is a subset of A^*

Examples

$$A = \{0, 1\}$$

$$\text{so } A^* = \{ "", "0", "1", "00", "01", "10", "11", "000", \dots \}$$

The set of bytes is a finite language over A :

$$\text{Bytes} = \{ b_0 b_1 b_2 b_3 b_4 b_5 b_6 b_7 \mid b_i \in A \}$$

The set of even-length bitstreams is an infinite language over A :

$$\text{Evens} = \{ "00", "01", "10", "11" \}^*$$


Prop as a formal language

- Alphabet:
 - { "A", "B", "C", ..., "(", ")", ... }
- Tokens:
 - Keywords: AND, OR, NOT
 - Literals: TRUE, FALSE
 - Punctuation: L_PAREN, R_PAREN
 - Input names: all nonempty subsets of A^* containing only letters
- Expressions:
 - The subset of Tokens* corresponding to valid circuits

How do we specify these details?

Keywords in Prop


PropKW = "AND" | "OR" | "NOT"



Finite automata and regular languages

Terminology

A *finite automaton* (or *finite state machine*) describes a system that can *transition* (or *move*) between different *states* in response to particular *inputs*.



Finite automata building blocks




state, labeled n



start state



accept (final) state



transition on input a



transition without consuming any input

Determinism

- A machine is non-deterministic (an *NFA*, or *non-deterministic finite automaton*) if it has a state with either more than one transition on the same symbol or if it has any ϵ -transitions.
- Otherwise the machine is deterministic (a *DFA*, or *deterministic finite automaton*)

Regular languages

A regular language is a language specified by:

- A deterministic finite automaton
- A nondeterministic finite automaton
- A regular expression

These are all equivalent: a DFA can be converted to a regex, a regex can be converted to an NFA, etc.

Regular languages

- $\{\}$ and $\{""\}$ are languages over any alphabet
- if $c \in A$, then $\{c\}$ is a language over A
- if L_1 and L_2 are regular languages, then so are
 - $L_1 \cup L_2 = \{x \mid x \in L_1 \text{ or } x \in L_2\}$
 - $L_1 \cdot L_2$ (or just $L_1 L_2$) = $\{xy \mid x \in L_1 \text{ and } y \in L_2\}$
- If L is a language, then so is
 - $L^* = \{""\} \cup \{xy \mid x \in L, y \in L^*\}$

Regular expressions

- A language for describing regular languages
- Often used in:
 - text editors/programming languages, for describing patterns in text strings (e.g. email address validation)
 - lexers, for describing the lexical structure of a language

Regular expressions

- ϵ

Empty: matches the empty string

- c

Constant: matches the single character 'c'

- $r_1 \mid r_2$

Alternation: strings matching either r_1 or r_2

- $r_1 \cdot r_2$ (or just $r_1 r_2$)

Sequencing/concatenation: a string matching r_1 followed by a string matching r_2

- r^*

Repetition: a sequence of zero or more strings, each matching r

- (r)

Grouping: string matching r

Derived forms

- $r^+ = r \cdot r^*$
 - Repetition: a sequence of one or more strings, each matching r
- $r? = \varepsilon \mid r$
 - Option: an empty string or a string matching r
- $[abc] = a \mid b \mid c$
 - Character class: any listed character (also allows ranges, e.g. $[a-zA-Z]$)
- $.$
 - Wildcard: matches any character
- $^$
 - Line start: matches the empty string, but only at the start of a line
- $\$$
 - Line end: matches the empty string, but only at the end of a line


Compiler-related examples

- Decimal integer literals: `[0-9]+`
- Keywords: `i • f, e • l • s • e`
- Haskell variables: `[a-z] [A-Za-z0-9' _]*`
- Haskell types/constructors: `[A-Z] [A-Za-z0-9' _]*`
- Whitespace: `[\t\n]*`
- C-style comment: `//.*$`

Abstract syntax of regexes

How do we write programs that operate on regexes?

- The same way we operate on any kind of syntax: with an abstract syntax data structure



Denotational semantics of regexes

Each regular expression describes a regular language, where $L(r)$ is the language denoted by r :

- $L(\varepsilon) = \{""\}$
- $L(c) = \{c\}$
- $L(r_1 \mid r_2) = L(r_1) \cup L(r_2) = \{x \mid x \in L(r_1) \text{ or } x \in L(r_2)\}$
- $L(r_1 \cdot r_2) = L(r_1) \cdot L(r_2) = \{xy \mid x \in L(r_1) \text{ and } y \in L(r_2)\}$
- $L(r^*) = L(r)^* = \{""\} \cup \{xy \mid x \in L(r), y \in L(r^*)\}$
- $L((r)) = L(r)$

This function is an interpreter, mapping a regex (syntax) to a set of strings (semantics).

Basics of lexical analysis

- Lexical analysis is carried out by a lexer/scanner/tokenizer
- Goal: to recognize and identify the sequence of tokens represented by the characters in a program's text
- The lexical structure (definition of tokens) is an important part of many language specifications

Lexemes

- A *lexeme* is a string that might represent a single atomic syntactic unit
- Examples of lexemes in Haskell:
 - `"0.0"`
 - `"String"`
 - `"True"`
 - `"if"`
 - `"("`
 - `"eval"`

Tokens

- A *token type* classifies lexemes; a *token* is a lexeme tagged with a token type
- Examples of tokens in Haskell:
 - `"0.0"` = `NUM(0.0)`
 - `"String"` = `ID("String")`
 - `"True"` = `ID("True")`
 - `"if"` = `IF`
 - `"("` = `L_PAREN`
 - `"eval"` = `ID("eval")`
- When a token type contains only one lexeme (e.g. `IF`), we usually leave out the lexeme and just write the type
- The tokens and lexemes for a language are usually chosen so that each valid lexeme is a member of exactly one token set

Patterns

- A *pattern* is a description of the way that a set of lexemes are written
- Informally, in natural language
 - e.g. from the Java spec: "An identifier is an unlimited-length sequence of Java letters and Java digits..."
- Formally, in the language of regular expressions:
 - `ID = letter • (letter | digit)*`

Common token types

- Keywords, symbols, punctuation
 - `for, if, then, <=, +, (, ;, .`
- Literals/constants
 - integers
 - floating point numbers
 - characters
 - strings
- Identifiers
 - `String, True, eval`

Other input elements

Other elements that might appear in the input stream (but are not tokens):

- Whitespace (space, tab, newline, etc.)
 - Except in whitespace-sensitive languages (e.g. Python)
- Comments

These are filtered out during lexing and not passed as tokens to the parser.

Lexical analysis summary

- Lexing breaks input streams of characters into output streams of tokens, usually filtering out whitespace and comments
- Regular expressions, regular languages, and finite automata provide a solid (but not mandatory) foundation for lexical analysis
 - Precise and concise notions for describing syntax
 - Expressive enough for the lexical syntax of many languages
 - Algorithms and practical tools exist to construct efficient lexers

Context-free languages

Matching brackets

- $\text{Brackets} = \{ " " \} \cup \{ [b] \mid b \in \text{Brackets} \}$
- So the words in Brackets are:
 - $"", "[]", "[[]]", "[[[]]]", "[[[[]]]]", \dots$
- In other words, nested pairs of bracket characters:
 - A sequence of N open brackets followed by N close brackets
- A subset of any language that uses brackets
- Is it regular?
 - Is there a regular expression r such that $L(r) = \text{Brackets}$?

Brackets is not regular

Remember the pumping lemma?

- In short, intuitively: DFAs can't count
- If s_n is the state that we reach after n open brackets and $n \neq m$, then $s_n \neq s_m$
- So there needs to be one state for each possible number of open brackets
- There are an infinite possible numbers of open brackets!
- So any machine to match Brackets must have infinite states, and therefore is not a **finite** state machine
- So Brackets can't be regular

Repetition vs. recursion

- Regular expressions allow iteration (repetition) (e.g. with the $*$ operator) but don't allow recursion (self-reference)
- A recursive characterization of Brackets is straightforward:

$$B \rightarrow \varepsilon$$

$$B \rightarrow [B]$$

- Meaning: an element of Brackets is either the empty string, or an element of Brackets surrounded by brackets

Context-free grammars

Formally: a context-free grammar $G = (T, N, P, S)$ consists of

- A set T of *terminal* symbols (tokens)
- A set N of *nonterminal* symbols
- A set P of *productions* (elements of $N \times (T \cup N)^*$)
 - Usually written " $n \rightarrow w$ " where $n \in N$ and $w \in (T \cup N)^*$
- A *start symbol* $S \in N$

Brackets CFG

$$T = \{ '[', ']' \}$$

$$N = \{ B \}$$

$$P = \{ B \rightarrow \varepsilon, B \rightarrow [B] \}$$

$$S = B$$

$$\text{Brackets} = (T, N, P, S)$$

In practice, we usually just write the productions; the start symbol is either denoted with S or assumed to be the left-hand symbol in the first production.

Example CFGs

Prop (without inputs) =

$P \rightarrow \text{TRUE}$

$P \rightarrow \text{FALSE}$

$P \rightarrow (P)$

$P \rightarrow \text{AND } P \ P$

$P \rightarrow \text{OR } P \ P$

$P \rightarrow \text{NOT } P$

Regex =

$R \rightarrow c$

$R \rightarrow \varepsilon$

$R \rightarrow R \cdot R$

$R \rightarrow R \mid R$

$R \rightarrow R^*$

Arithmetic =

$E \rightarrow n$

$E \rightarrow (E)$

$E \rightarrow E + E$

$E \rightarrow E * E$

Derivations

Formally: A *derivation* of a CFG is a sequence of strings $s_1 \rightarrow s_2 \rightarrow \dots \rightarrow s_n$ where each string s_{i+1} is obtained from the previous string s_i by choosing a production $n \rightarrow w$ and replacing an occurrence of n in s_i with w .

In Brackets:

$B \rightarrow [B] \rightarrow [[B]] \rightarrow [[[B]]] \rightarrow [[[]]]$

In Prop:

$P \rightarrow \text{AND } P \ P \rightarrow \text{AND } (P) \ P \rightarrow \text{AND } (\text{NOT } P) \ P \rightarrow \text{AND } (\text{NOT TRUE}) \ P \rightarrow \text{AND } (\text{NOT TRUE}) \text{ FALSE}$

We say that a CFG *generates* the language that contains all strings that can be derived from the start symbol; any language generated by a CFG is a *context-free language*.

Why "context-free"?

- The productions in a CFG can be expanded anywhere in a derivation, regardless of surrounding symbols
- In contrast to a context-sensitive grammar, which can have productions with the left hand side restricted to certain contexts - e.g. $[B] \rightarrow (B)$

EBNF grammars

Extended Backus-Naur form (EBNF) is a shorthand syntax for CFGs, very often used in programming language specifications.

Definition: ... = ...

Terminal string: "..."

Alternation: ... | ...

Zero or one: [...]

Zero or more: { ... }

```
digit = "0" | "1" | "2" | ... | "9"
```

```
expr = [digit]{digit} {"+" expr}
```

Parse trees

Multiple derivations

Often, there are multiple choices for a step in a derivation:

- In a *right-most* derivation, replace the right-most nonterminal at each step
- In a *left-most* derivation, replace the left-most nonterminal at each step
- Any other arbitrary ordering

Does it matter which derivation we use?

Example: $1 + 2 * 3$

Arithmetic =

$E \rightarrow n$

$E \rightarrow (E)$

$E \rightarrow E + E$

$E \rightarrow E * E$

Left-most:




$$\begin{aligned} & E \\ \rightarrow & E + E \\ \rightarrow & 1 + E \\ \rightarrow & 1 + E * E \\ \rightarrow & 1 + 2 * E \\ \rightarrow & 1 + 2 * 3 \end{aligned}$$

Right-most:

$$\begin{aligned} & E \\ \rightarrow & E + E \\ \rightarrow & E + E * E \\ \rightarrow & E + E * 3 \\ \rightarrow & E + 2 * 3 \\ \rightarrow & 1 + 2 * 3 \end{aligned}$$

Parse trees

To capture the structure of a derivation, we use a graphical tree notation:



- These are called *parse trees*, or sometimes *concrete syntax trees (CSTs)*
 - More information than an AST (e.g. parens)
- In theory, token stream \rightarrow parse tree \rightarrow AST
- In practice, parse trees are often left implicit (token stream \rightarrow AST)
 - A parse tree is the call graph of a recursive descent parser

Arithmetic =

$E \rightarrow n$

$E \rightarrow (E)$


$E \rightarrow E + E$

$E \rightarrow E * E$

Example: $1 + 2 * 3$

Leftmost:

E
 $\rightarrow E + E$
 $\rightarrow 1 + E$
 $\rightarrow 1 + E * E$
 $\rightarrow 1 + 2 * E$
 $\rightarrow 1 + 2 * 3$



Rightmost:

E
 $\rightarrow E + E$
 $\rightarrow E + E * E$
 $\rightarrow E + E * 3$
 $\rightarrow E + 2 * 3$
 $\rightarrow 1 + 2 * 3$

Both derivation orders produce the same parse tree - the only difference is the order in which the nodes are constructed.

Arithmetic =

$E \rightarrow n$

$E \rightarrow (E)$


$E \rightarrow E + E$

$E \rightarrow E * E$

Example: $1 + 2 + 3$


Leftmost:

E
 $\rightarrow E + E$
 $\rightarrow E + E + E$
 $\rightarrow 1 + E + E$
 $\rightarrow 1 + 2 + E$
 $\rightarrow 1 + 2 + 3$



Rightmost:

E
 $\rightarrow E + E$
 $\rightarrow E + E + E$
 $\rightarrow E + E + 3$
 $\rightarrow E + 2 + 3$
 $\rightarrow 1 + 2 + 3$



The two orders produce different trees!

Ambiguity

- A grammar is *ambiguous* if the language it generates contains a string with more than one parse tree
- e.g. our simple arithmetic grammar is ambiguous because "1+2+3" has multiple parse trees
 - There are many other expressions in the language with multiple parse trees, but one is enough to demonstrate ambiguity
- Ambiguity is a property of a grammar, not a language
 - We can have multiple grammars describing the same language, some ambiguous and some unambiguous

Dealing with ambiguity

Does it matter?

- If all parse trees for a string are semantically equivalent, it doesn't
 - e.g. for regexes, $r_1 \cdot (r_2 \cdot r_3)$ describes exactly the same language as $(r_1 \cdot r_2) \cdot r_3$, so we can parse $r_1 \cdot r_2 \cdot r_3$ either way arbitrarily without issue
- If different trees have different meanings, we need to choose between them
 - Disambiguating rules (e.g. operator precedence)
 - Rewrite the grammar to avoid ambiguity

Precedence and associativity

For two arbitrary infix operators, \boxplus and \oplus :

- If \boxplus has *higher precedence* than \oplus , then " $a \boxplus b \oplus c$ " parses as " $(a \boxplus b) \oplus c$ "
- If \boxplus is *left-associative*, then " $a \boxplus b \boxplus c$ " parses as " $(a \boxplus b) \boxplus c$ "
- If \boxplus is *right-associative*, then " $a \boxplus b \boxplus c$ " parses as " $a \boxplus (b \boxplus c)$ "
- If \boxplus is *non-associative*, then " $a \boxplus b \boxplus c$ " is a syntax error

Fixity = precedence + associativity

Order of operations

- There are widely used conventions for the precedence of arithmetic operators (PEMDAS)
- What about less traditional operators?
 - Ternary conditionals ($x \text{ ? } y \text{ : } z$) in C/C++/Java
 - User-defined operators ($.@.$) in Haskell
- Rules vary by language
 - C/C++/Java have a table in the spec with the fixity of each operator
 - Haskell allows user-specified fixity (e.g. `infixl 2 (.@.)`)
 - `infix`: non-associative
 - `infixl`: left-associative
 - `infixr`: right-associative

An unambiguous grammar for expressions

- Three nonterminals:
 - **E**xpressions: sum of products
 - **P**roducts: product of atoms
 - **A**toms: parenthesized expressions and numbers
- * has higher precedence than +, and both associate to the left
- Choose (arbitrarily) to only allow left-most derivations

Arithmetic =

$E \rightarrow P$

$E \rightarrow E + P$

$P \rightarrow A$

$P \rightarrow P * A$

$A \rightarrow (E)$

$A \rightarrow n$

Example: $1 + 2 * 3$

E
→ **E** + P
→ **P** + P
→ **A** + P
→ 1 + **P**
→ 1 + **P** * A
→ 1 + **A** * A
→ 1 + 2 * **A**
→ 1 + 2 * 3

Arithmetic =

$E \rightarrow P$

$E \rightarrow E + P$

$P \rightarrow A$

$P \rightarrow P * A$

$A \rightarrow (E)$

$A \rightarrow n$

Context-free languages summary

- Context-free grammars describe a significantly larger family of languages than regular expressions
 - Including most programming languages
- Parse trees are graphical descriptions of CFG derivations
 - Reflect the grammatical structure of the input
 - Highlight ambiguities in the grammar
 - Include more detail than ASTs
- Operator precedence and associativity can reduce/eliminate ambiguity